

Deep Stock Representation Learning: From Candlestick Charts to Investment Decisions

Guosheng Hu^{*6}, Yuxin Hu^{*1}, Kai Yang², Zehao Yu³, Flood Sung⁴, Zhihong Zhang³
Fei Xie¹, Jianguo Liu¹, Neil Robertson⁶, Timothy Hospedales⁷, Qiangwei Miemie^{5,8}

¹Shanghai University of Finance and Economics ²University of Shanghai for Science and Technology ³Xiamen University ⁴Independent Researcher
⁵Queen Mary University of London ⁶Queen's University Belfast ⁷The University of Edinburgh ⁸Yang's Accounting Consultancy Ltd
{huguosheng100, floodsung}@gmail.com, {hyx552211, yang_kai_2008}@163.com, 1127329344@qq.com, zhihong@xmu.edu.cn
xiefei@mail.shufe.edu.cn, liujg004@ustc.edu.cn, n.robertson@qub.ac.uk, t.hospedales@ed.ac.uk, yongxin@yang.ac

Abstract

We propose a novel investment decision strategy based on deep learning. Many conventional algorithmic strategies are based on raw time-series analysis of historical prices. In contrast many human traders make decisions based on visually observing candlestick charts of prices. Our key idea is to endow an algorithmic strategy with the ability to make decisions with a similar kind of *visual* cues used by human traders. To this end we apply Convolutional AutoEncoder (CAE) to learn an asset representation based on visual inspection of the asset's trading history. Based on this representation we propose a novel portfolio construction strategy by: (i) using the deep learned representation and modularity optimisation to cluster stocks and identify diverse sectors, (ii) picking stocks within each cluster according to their Sharpe ratio (Sharpe 1994). Overall this strategy provides low-risk high-return portfolios. We use the Financial Times Stock Exchange 100 Index (FTSE 100) data for evaluation. Results show our portfolio outperforms FTSE 100 index and many well known funds in terms of total return in 2000 trading days.

Introduction

Investment decision making is a classic research area in quantitative and behavioural finance. One of the most important decision problems is portfolio construction and optimisation (Markowitz 1952; Kelly 1956), which addresses selection and weighting of assets to be held in a portfolio. Financial institutions try to construct and optimise portfolios in order to maximise investor returns while minimising investor risk.

It is hard to predict the stock market due to partial information and the involvement of irrational traders (Fama 1998) (who tend to over- or under-react on stock price (Barberis, Shleifer, and Vishny 1998)). Nevertheless, behavioural finance argues the past price record impacts the market future performance (Balsara and Zheng 2006). Therefore many portfolio selection strategies are built on 'judging and following' historical data. Existing algorithmic methods to construct portfolios use raw time series data as input to drive optimisation via time-series analysis methods. However, human traders typically observe price history visually, for example in the form of candlestick charts in order

to make trading decisions. Candlestick charts (Fig. 2) provide a visual representation of four price parameters: lowest, highest, opening and closing price for the day. Thus current algorithmic strategies interpret the underlying data in a very different way than human traders due to the gap between raw time series and visual representations such as charts.

In this study we aim to endow an algorithmic portfolio optimisation method with the ability to make decisions based on price history interpreted in a visual way more similar to human traders. In particular we develop a deep learning (DL) approach to representing price history visually and eventually driving portfolio construction. Convolutional DL approaches such as Convolutional AutoEncoder (CAE, unsupervised learning) (Masci et al. 2011) and Convolutional Neural Network (CNN, supervised learning) (LeCun et al. 1998), have achieved very impressive performance for analysing visual imagery. This has motivated researchers to adapt CNNs and CAEs to applications that are not naturally image-based. It is typically achieved by converting raw input signals from other modalities into images to be processed by CNNs or CAEs. Good results have been achieved for diverse applications in this way because of DL's efficacy in capturing non-linearity, translation invariance, and spatial correlations. For example, traditional speech recognition methods used the 1-D signal vector, i.e. the raw input waveform, or its frequency domain projection (Rabiner and Juang 1993; Povey et al. 2011). In contrast an alternative approach is to convert the 1-D signal to a spectrogram, i.e. an image, in order to leverage the strength of CNNs to achieve promising recognition performance (Amodi et al. 2016). As another well known example, AlphaGo (Silver et al. 2016) represents the board position as a 19×19 image, which is fed into a CNN for feature learning. With similar motivation, we explore converting raw input data, (i.e., time-series) to images in order to leverage DL for better stock representation learning, and hence ultimately better portfolio optimisation.

The first novelty of this study is exploiting deep learning to interpret price history data in a way more similar to human traders by 'seeing' candlestick charts. To fully evaluate our approach, we further construct a complete portfolio generation pipeline which requires a number of other elements. Our full investment strategy includes: (1) deep feature learning by visual interpretation price history, (2) clustering the deep representation in order to provide a data-driven seg-

* These authors contributed equally to this work

mentation of the market, (ii) visual portfolio construction.

For *visual representation* learning we take a 4-channel stock time-series and synthesise candlestick charts to present price history as images. As we generate millions of training images, we choose unsupervised CAE for representation learning, avoiding expensive big data labelling. The generated charts are fed into a deep CAE to learn representations which effectively encode the semantics of the historical time-series. In the next *clustering* step, we aim to segment the market into diverse sectors in a data-driven way based on our generated visual representation. This is important to provide risk reduction by selecting a well diversified portfolio (Nanda, Mahanty, and Tiwari 2010; Tola et al. 2008). Popular clustering methods such as K-means are not suitable here because they are non-deterministic and/or require a pre-defined numbers of clusters to find. In particular non-deterministic methods are not acceptable to real financial users. To address this we adapt the modularity optimization method (Newman 2006) – originally designed for network community structure – to stock clustering. Specifically, we construct a network (graph) of stocks where each node corresponds to a stock, and edges represent the correlation between a pair of stocks (as represented by the learned visual feature vectors). To complete the pipeline we perform *portfolio construction* by the simple but effective approach of choosing the best stock within each cluster according to their Sharpe ratio (Sharpe 1994). As we will see in our evaluation, this strategy provides a surprisingly approach to constructing portfolios that combine high returns with low-risk.

Our contributions can be summarised as follows:

- A novel deep learning approach that imitates human traders' ability to 'see' candlestick charts. We convert raw 4-D stock time-series to candlestick images and use unsupervised convolutional autoencoders to learn a representation for price history that encodes semantic information about the 4-D time series. To our knowledge, we are the first to apply deep learning techniques from computer vision to investment decisions.
- To support low-risk diversified portfolio construction, we segment the market by clustering stocks based on their learned deep features. To achieve deterministic clustering without needing to specify the number of clusters, we adapt modularity optimisation, originally used for detecting community structure in networks, to stock clustering.
- We evaluate our strategy using Financial Times Stock Exchange 100 Index (FTSE 100) stock data. Results show (a) our learned deep stock representation effectively captures semantic information. (b) our portfolio construction strategy outperforms the FTSE 100 index and many well-known funds including 2 big funds (CCA, VXX) and the 3 best performed funds (IEO, PXE, PXI) recommended by Yahoo in terms of total return in 2000 trading days, showing the effectiveness of our investment strategy.

Related Work

Investment Decision Strategy Portfolio optimisation is a mainstay of investment decision making. There are two

mainstream directions of investing: (i) determining the optimal allocation of capital with different perspectives on period selection: (i) The Mean Variance Theory (Markowitz 1952; Markowitz 1959; Markowitz and Todd 2000) chooses a single-period and optimises the portfolio by the trading-off of the expected return and the variance of return. Derived from this strategy, many well-known ratios were proposed as objective functions to drive portfolio optimisation. For example the Sharpe ratio (Sharpe 1994) and Omega ratio (Keating and Shadwick 2002). In other words, the mean variance theory directs investment based on the 'risk-return' profile of a portfolio. (ii) The Capital Growth Theory (Kelly 1956; Hakansson and Ziemba 1995) defines the optimal strategy in dynamic and reinvestment situations and primarily focuses on the growth rate. Unlike mean variance theory, capital growth pays great attention to time dependence. It analyses the information of different time series, such as asset trees (Onnela et al. 2003), which uses time dependence to construct a tree based on the best classification of stocks. Our strategy combines the ideas of (i) and (ii) by using Sharpe ratio and encoding time dependence with deep learning.

Deep Feature Learning Effective visual representations 'features' are essential for performing computer vision tasks effectively. Hand-crafted features such as SIFT (Lowe 1999) and their aggregation mechanisms such as Fisher vectors (Perronnin, Sánchez, and Mensink 2010) have been the workhorse of computer vision for decades. However they ultimately suffer from a performance bottleneck because downstream errors cannot be fed back to improve their responsiveness to relevant observations and their robustness to nuisance factors. Unlike hand-crafted features deep learning approaches exploit end-to-end learning by back propagation of downstream evaluation errors to improve the efficacy of the learned features themselves. Deep feature learning can easily be categorised into unsupervised deep feature learning (UDFL) and supervised deep feature learning (SDFL). UDFL methods include deep belief networks (DBNs) and AutoEncoders (AEs) and their variants. (Convolutional) deep belief network achieved great success on both computer vision (Lee et al. 2009a) and speech processing (Lee et al. 2009b) tasks. The AE and its variants such as the denoising AE (Vincent et al. 2008) and variational AE (Kingma and Welling 2013) also show strong feature learning capacity in many tasks. For SDFL, CNNs (LeCun et al. 1998) are most widely applied. Since AlexNet (Krizhevsky, Sutskever, and Hinton 2012) achieved record breaking object recognition performance, CNNs have been widely used in many computer vision tasks. Since AlexNet, researchers have proposed a stream of increasingly effective CNN architectures such as Inception (Szegedy et al. 2015), VGGNet (Simonyan and Zisserman 2014), ResNet (He et al. 2016) and DenseNet (Huang et al. 2016). Compared with UDFL, SDFL can generally learn better features because of the labeled training data. However, it is a nontrivial task to label big data. In this work we will need to programatically generate millions of candlestick images, for which data labelling is not feasible. Therefore we use unsupervised convolutional autoencoders (CAEs) for visual feature learning.

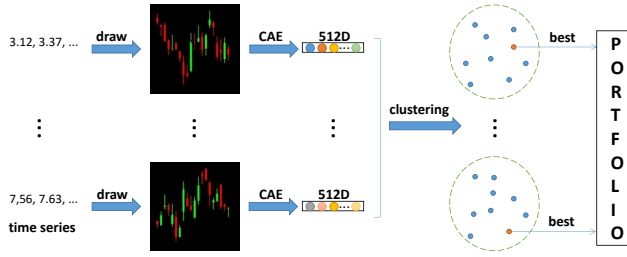


Figure 1: Schematic illustration of our investment decision pipeline. The architecture of the CAE is detailed in Fig. 3.

Methodology

Our overall investment decision pipeline includes three main modules: deep feature learning, clustering, and portfolio construction. For *deep feature learning*, raw 4-channel time series data describing stock price history are converted to standard candlestick charts. These charts are then fed into a deep CAEs for visual feature learning. These learned features provide a vector embedding of a historical time-series that captures key quantitative and semantic information. Next we *cluster* the features in order to provide a data-driven segmentation of the market to underpin subsequent selection of a diverse portfolio. Many common clustering methods are not suitable here because they are non-deterministic or require predefinition of the number of clusters. Thus we adapt modularity optimisation for this purpose. Finally, we perform *portfolio construction* by choosing stocks with the best performance measured by Sharpe ratio (Sharpe 1994) from each cluster. The overall pipeline is summarised schematically in Fig. 1. Each component is discussed in more detail in the following sections.

Autoencoders

An AutoEncoder (AE) is a typical artificial neural network for unsupervised feature learning. It learns an effective representation without requiring data labels, by encoding and reconstructing the input via bottleneck. An AE consists of an encoder and a decoder: The encoder uses raw data as input and produces representation as output and the decoder uses the learned feature from the encoder as input and reconstructs the original input as output.

For simplicity, we describe formally only a fully-connected AE with 3 layers: encoding, hidden and decoding layer. The encoder f projects the input \mathbf{x} to the hidden layer representation (encoding) \mathbf{h} . Typically, f includes a linear transform followed by a nonlinear one:

$$\mathbf{h} = f(\mathbf{x}) = \sigma(W\mathbf{x} + \mathbf{b}) \quad (1)$$

where W is the linear transform and \mathbf{b} is the bias. σ is the nonlinear (activation) function. There are many ‘activation’ functions available such as sigmoid, tanh and RELU. We use RELU (REctified Linear Units) in this work.

The decoder g maps the hidden representation \mathbf{h} back to the input \mathbf{x} via another linear transform W' and the same ‘activation’ function σ as shown in Eq. (2):

$$\mathbf{x} = g(\mathbf{h}) = \sigma(W'\mathbf{h} + \mathbf{b}') \quad (2)$$



Figure 2: Stock candlestick chart. (a) Candlestick explanation (b) One sample chart from stock ‘Royal Dutch Shell Plc Class A’ between 03/01/2017 and 31/01/2017.

To optimise $\{W', W, \mathbf{b}, \mathbf{b}'\}$, the AE uses the cost function:

$$\min_{W, \mathbf{b}, W', \mathbf{b}'} \sum_{i=1}^N \|\mathbf{x}_i - g(f(\mathbf{x}_i))\|^2 \quad (3)$$

where \mathbf{x}_i represents the i th one of N training samples. Using the objective in Eq. (3) the AE parameters can be learned via mini-batch stochastic gradient descent.

Convolutional Autoencoders Traditional AEs ignore the 2D image structure. Convolution AutoEncoders (CAEs) (Masci et al. 2011) extend vanilla AEs with convolution layers that help to capture the spatial information in images and reduce parameter redundancy compared to AEs. We use convolutional autoencoders in this work to process stock charts visually.

Deep Feature Learning with Convolutional Autoencoders

Chart Encoding To realise an algorithmic portfolio construction method based on visual interpretation of stock charts, we need to convert price history data to an image representation. Our raw data for each stock is a 4-channel time series (the lowest, the highest, open, and closing price for the day) in a 20-day time sequence. We use computer graphics techniques to convert these to a candlestick chart represented as a RGB image as shown in Fig. 2. The whisker plots describe the four raw channels, with color coding describing whether the stock closed higher (green) or lower (red) than opening. An encoded candlestick chart image provides the visual representation of one stock over a 20-day window for subsequent visual interpretation by our deep learning method.

Convolutional Autoencoder Our CAE architecture is summarised in Fig. 3. It is based on the landmark VGG network (Simonyan and Zisserman 2014), specifically VGG16. The VGG network is a highly successful architecture initially proposed for visual recognition. To adapt it for use as a CAE encoder, we remove the final 4096D FC layers from VGG-16 and replace them by an average pooling layer to generate one 512D feature. The decoder is a 7-layer deconvolutional network that starts with a 784D layer that is




Figure 3: Overview of our CAE architecture. We follow the encoder (top) - decoder (bottom) framework. The 512D feature following average pooling provides our representation for clustering and portfolio construction.

fully connected with the 512D embedding layer. Following 6 upsampling deconvolution layers eventually reconstruct the input based on our 512D feature. When trained with a reconstruction objective, the CAE network learns to compress input images to the 512D bottleneck in a manner that preserves as much information as possible in order to be able to reconstruct the input. Thus this single 512D vector encodes the 20-day 4-channel price history of the stock, and will provide the representation for further processing (clustering and portfolio construction).

Clustering

We next aim to provide a mechanism for diversified – and hence low risk – portfolio selection. Unlike some existing methods which calculate the stock correlation based on the raw time series, in this paper, we use our learned features in a clustering step in order to segment the market. As discussed before many existing clustering methods are non-deterministic or require pre-specification of the number of clusters, which make them unsuited for our application domain. To solve these problems, we introduce the network modularity method to find the cluster structure of the stocks, where each stock is set as one node and the link between each pair of stocks is set as the correlation calculated by the learned features (Blondel et al. 2008). Modularity is the fraction of the links that fall within the given group minus the expected fraction if links are distributed at random. Modularity optimisation is often used for detecting community structure in networks. It operates on a graph (of stocks in our case), and updates the graph to group stocks so as to eventually achieve maximum modularity before terminating. Thus it does not need a specified number of clusters and is not affected by initial node selection, which

Algorithm Description One 20-day history of the entire market is represented as a single graph. Each stock corresponds a single node in the initial graph, and is represented by our CAE learned feature. The cosine similarity between these features is then used to weight the edges of the graph connecting each stock according to their similarity.

The Blondel method (Blondel et al. 2008) for cluster de-

tection provides a fast greedy approach based on modularity optimisation (Newman 2006). It can be divided into two iterative steps:

1. Initially, each node in network is considered a cluster. Then, for each node i , one considers its neighbour j and calculates the modularity increment by moving i from its cluster to the cluster of j . The node i is placed in the cluster where this increment is maximum, but only if this increment is positive. If no positive increment is possible, the node i stays in its original cluster. This process is applied repeatedly for all nodes until no further increment can be achieved.
2. The second step is to build a new network whose nodes are the clusters found in the first step. The weights of the links between the new nodes are given by the sum of the number of the links between nodes from the corresponding component clusters now included in that node.

The key notion of modularity increment $\Delta Q(x, y)$, where x is one of the new nodes and y is x 's neighbour, is defined as:

$$\Delta Q(x, y) = \left[\frac{e_{in, c_y} + 2k_{x, c_y}}{2m} - \left(\frac{e_{tot, c_y} + 2k_x}{2m} \right)^2 \right] - \left[\frac{e_{in, c_y}}{2m} - \left(\frac{e_{tot, c_y}}{2m} \right)^2 - \left(\frac{k_x}{2m} \right)^2 \right], \quad (4)$$

where c_y is the cluster that contains the node y . e_{in, c_y} is the number of the links inside cluster c_y and e_{tot, c_y} denotes the sum of degrees of nodes in the cluster c_y , k_x is the degree of the node x , k_{x, c_y} is the sum of the links from the node x to the nodes in the cluster c_y and m is the number of all the links in the network.

The end result is a segmentation of the market into a maximum modularity graph with K nodes/clusters, where each node contains all the stocks that cluster together. It should be noticed that whether the graph is fully connected or not, the Blondel method finds a hierarchical cluster structure.

Portfolio Construction and Backtesting

So far we have described the parts of our strategy designed to maintain low risk via high diversity (stock clustering driven by learned deep features). Given the learned stock clustering (market segmentation), we finally describe how a complete portfolio is constructed by picking diverse yet high-return stocks, as well as our portfolio evaluation process.

Stock Performance Return (profit on an investment) is defined as $r_t = (V_f - V_i)/V_i$, where V_f and V_i are the final and initial values, respectively. For example, to compute daily stock return, V_f and V_i are closing prices of today and yesterday, respectively. We measure the performance of one particular stock over a period using the Sharpe ratio (Sharpe 1994) $s = \bar{r}/\sigma_r$, where \bar{r} is the mean return, σ_r is the standard deviation over that period. Thus the Sharpe ratio s encodes a tradeoff of return and stability. Maximum Drawdown (MDD) is the measure of decline from peak during a specific period of investment: $MDD = (V_t - V_p)/V_p$, where V_t and V_p mean the trough and peak values, respectively.

Training and Testing For every 20 trading days, we cluster all the stocks as described in the previous section. To actually construct a portfolio we then choose the stock with



the highest Sharpe ratio (Sharpe 1994) within each cluster. We then hold the selected portfolio for 10 days. Over these following 10 days, we evaluate the portfolio by computing our ‘compound return’ for each selected stock. The overall return of one portfolio in these 10 days is the average compound return of all the selected stocks. The stride for the 20-day window is 10 in this work.

The process of portfolio selection/optimisation in Finance is analogous to the process of training in machine learning, and the subsequent return computation is analogous to the testing process.

Since the modular-optimisation based clustering discovers the number of stocks in a data driven way, in different trading periods we may have different number of clusters. Assume that we obtain K_1 clusters in one period and will select K_2 stocks to construct one portfolio. Then, letting Q and R indicate quotient and remainder respectively in $[Q, R] = K_2/K_1$: K_2 stocks are picked by taking (i) Q stocks from each of the K_1 clusters and (ii) the remaining R best performing stocks across all K_1 clusters. Then, we allocate equally $1/K_2$ of the fund to each of the chosen stocks.

Experiments

We first introduce our dataset and experimental settings. Then we analyse the outputs of feature learning and clustering. Finally, we compare our investment strategy (feature extraction, clustering, portfolio optimisation) as a whole with other strategies.

Dataset and Settings



Preprocessing For evaluation, we use the stock data of Financial Times Stock Exchange 100 Index (FTSE 100). The FTSE 100 is a share index of the 100 companies listed on the London Stock Exchange with the highest market capitalisation. It effectively represents the economic conditions of the UK and even the EU. We use all the stocks in FTSE 100 from 4th Jan 2000 to 14th May 2017. We use the adjusted stock price (accounting for stock splits, dividends and distributions). Every 20-day 4-channel (the lowest price of the day, the highest price, open price, closing price) time series generates a standard candlestick chart. We generate 400K FTSE100 charts in all. The training images for our CAE are candlestick charts rendered as 224×224 images to suit our VGG16 architecture (Simonyan and Zisserman 2014).

Implementation Details The batch size is 64, learning rate is set to 0.001, and the learning rate decreases with a factor of 0.1 once the network converges.

Qualitative Results



Visualising and Understanding Deep Features First, we visualise the input and reconstructed candlestick charts in Fig. 4. We can see that the reconstructed charts effectively capture the variations of the inputs showing the efficacy of our CAE for unsupervised learning and the representational capacity of our 512D feature (CAE bottleneck layer). Note that as is commonly the case for autoencoders, the background is slightly grey because the pixel values of reconstructed background are close to 0, but not equal to 0.




Figure 4: Candlestick chart input (top) and reconstructed (bottom) by our CAE.




Figure 5: t-SNE visualisation of FTSE 100 CAE features. One color indicates one industrial sector. The stocks on the right are all from Sector Materials: RRS.L (Randgold Resources Limited), FRES.L (Fresnillo PLC), ANTO.L (Antofagasta PLC), BLT.L (BHP Billiton Ltd), AAL.L (Anglo American PLC), RIO.L (Rio Tinto Group), KAZ.L (KAZ Minerals), EVR.L (EVRAZ PLC)

Second, we visualise the learned features. The features of one whole year (2012) for one particular stock are concatenated to form a new feature. These new features of all the stocks are visualised in Fig. 5 using the t-Distributed Stochastic Neighbor Embedding (t-SNE) (Maaten and Hinton 2008) method. One color indicates one industrial sector defined by bloomberg (<https://www.bloomberg.com/quote/UKX:IND/members>). From Fig. 5, we can see the stocks with similar semantics (industrial sector) are represented close to each other in the learned feature space. For example, Materials related stocks are clustered. This illustrates the efficacy of our CAE and learned feature for capturing semantic information about stocks.

Analysis of Modularity Optimisation Clusters We next analyse the effect of our modularity optimisation based clustering based on our learned feature. Fig. 6 shows 3 clusters produced by our modularity optimisation-based clustering in terms of the return during both the 20-day training window for the clustering (above), and during the 10-day evaluation window (below). These graphs illustrate that the clusters represent different characteristic price trajectories. For

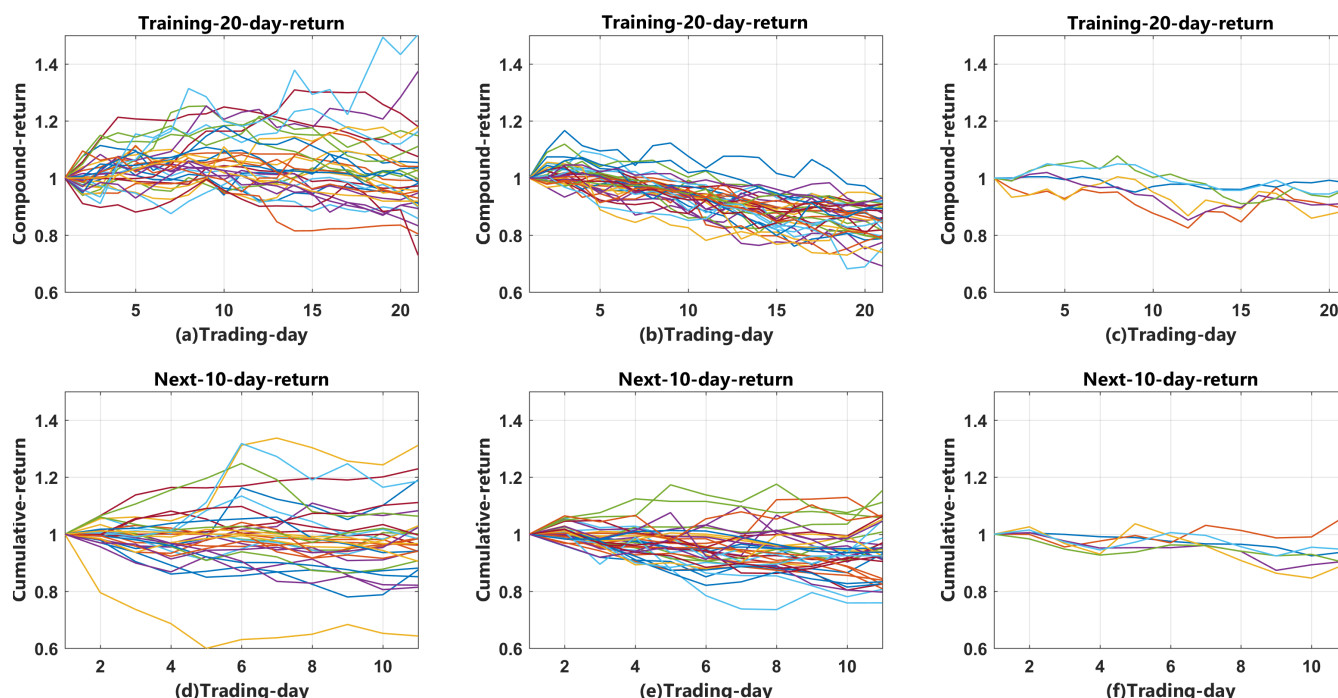


Figure 6: Visualisation of modularity optimisation clustering results. Row 1: compound returns of all the stocks in 3 clusters during portfolio optimisation (20 trading days), Row 2: compound returns during test (in the following 10 trading days).

example the top-left cluster represents a position of no net profit or loss, but high fluctuation. The top-middle cluster represents stocks with a characteristic of continuous steep decline; while the top-right cluster represents those with continuous slight decline. Clearly the segmentation method has allocated stocks to clusters with distinct trajectory characteristics. During testing the three corresponding clusters on the bottom behave quite similarly as those above. This demonstrates the momentum effect (Jegadeesh and Titman 1993), and the consistency between portfolio selection (20-day, training) and return computation (10-day, testing).

Quantitative Results

For quantitative evaluations, we apply 7 measures for evaluation: Total return, daily Sharpe ratio, max drawdown, daily mean return, monthly mean return, yearly mean return, and win year. Win year indicates the percent of the winning years. The other measures are defined in the section of ‘Portfolio Construction and Backtesting’. We choose $K_2 = 5$ stocks to construct all the portfolios compared.

Comparison with FTSE 100 Index We perform back-testing to compare our full portfolio optimisation strategy against the market benchmark (FTSE 100 Index). In Fig. 7, we compare with FTSE 100 index. Fig. 7 (a) shows the comparison over a long-term trading period (4K trading days, 31/01/2000 - 06/10/2016) showing the overall effectiveness of our strategy. Note that it is very difficult for funds to consistently outperform the market index over an extended period of time due to the complexity and diversity of market

variations. In Fig. 7 (b)-(c) we show specific shorter term periods where the market is behaving very differently including down-up (b), flat (c) and bullish (d). The overall dynamic trends of our strategy reflect the conditions of the market (meaning that the stocks selected by our strategy are representative of the market), yet we outperform the market even across a diverse range of conditions (b-c), and over a long time-period (a).

Feature and Clustering

An evaluation of features and clustering methods is also conducted over a long term period (4K trading days). From Table 2, the total return of our method (D-M, deep feature + modularity-based clustering) is higher than R-M (R-M, Raw time series + modularity-based clustering), 283.5% vs 208.8%. It means the deeply learned feature can capture richer information, which is more effective for portfolio optimisation than raw time series. Similar conclusions can be drawn based on other measures. In terms of clustering method, our modularity optimization method works better than D-K (deep feature + k -means) in terms of returns and daily Sharpe, showing the effectiveness of modularity-based clustering. As explained in the Introduction, k -means cannot be used for portfolio construction in practice. Specifically, the results of k -means cannot be repeated because of the randomness of the initial seed. Non-deterministic investment strategies are not acceptable to financial users in practice as they add another source of uncertainty (risk) that is hard to quantify.

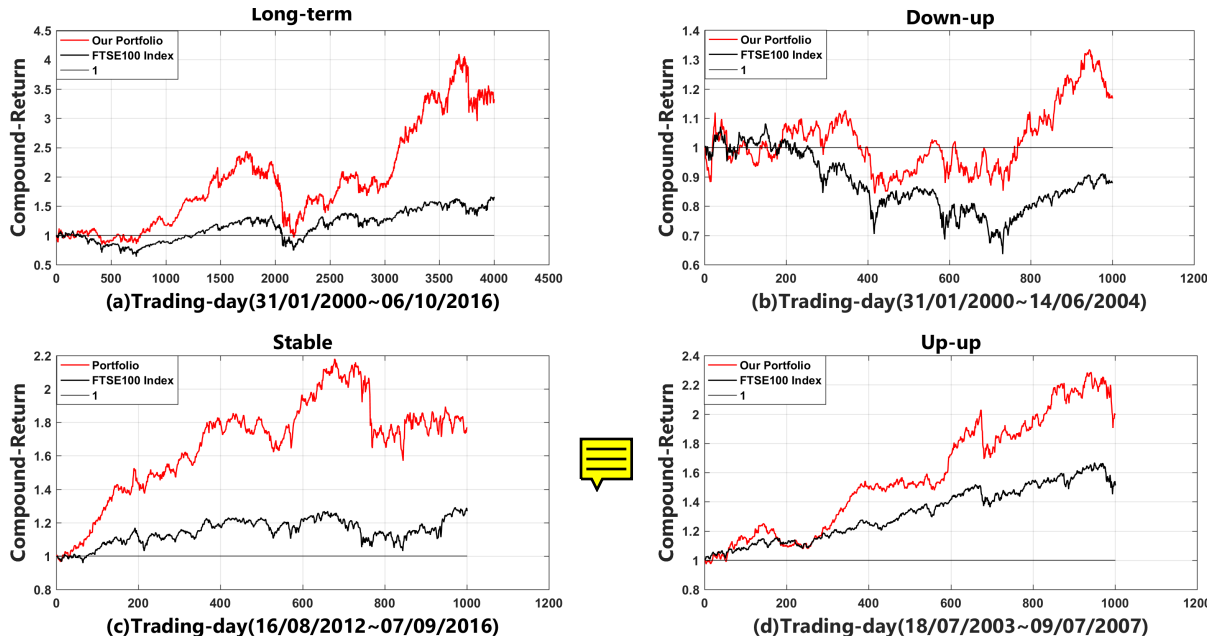


Figure 7: Our portfolio vs FTSE 100 Index

Table 1: Comparison with well-known Funds

	CCA	VXX	IEO	PXE	PXI	Ours
Total Return (↑)	117.0%	-99.9%	89.9%	101.6%	152.2%	215.4%
Daily Sharpe (↑)	0.7	-1.1	0.4	0.4	0.6	0.8
Max Drawdown (↑)	-22.2%	-99.9%	-56.8%	-57.6%	-59.3%	-30.9%
Daily Mean Return (↑)	10.9%	-67.7%	12.7%	13.4%	15.9%	16.7%
Monthly Mean Return (↑)	10.7%	-66.4%	11.5%	12.4%	14.9%	16.6%
Yearly Mean Return (↑)	6.5%	-44.6%	5.0%	8.2%	9.4%	11.8%
Win Years (↑)	62.5%	0.0%	62.5%	50.0%	75.0%	62.5%

Table 2: Comparison of Features and Clustering Methods.

	R-M	D-M (Ours)	D-K
Total Return (↑)	208.8%	283.5%	272.6%
Daily Sharpe (↑)	0.44	0.50	0.49
Max Drawdown (↑)	-55.6%	-60.5%	-59.0 %
Daily Mean Return (↑)	9.7 %	11.1 %	10.9%
Monthly Mean Return (↑)	8.7%	10.0 %	9.9%
Yearly Mean Return (↑)	9.6%	10.0 %	11.19%
Win Years (↑)	64.71%	69.52%	66.31%

Comparison with Funds To further analyse the effectiveness of our strategy, we compare our strategy with well known public funds in stock market. Specifically, we select 2 big funds (CCA and VXX) and the top 3 best performed funds (IEO, PXE, PXI) recommended by YAHOO (<https://finance.yahoo.com/etfs>). Note that the ranking of funds change over time. The fund data is obtained from Yahoo Finance. Because VXX starts from 20/01/2009, this evaluation is computed over 2K trading days (20/01/2009-09/01/2017). From Table 1, our portfolio achieved the highest returns: Total (215.4%), daily

(16.7%), monthly (16.6%), yearly (11.8%) in 2000 trading days, showing the strong profitability of our strategy. We also achieved the highest daily Sharpe ratio (0.8), meaning that we effectively balance the profitability and variance. We achieve the 2nd lowest max drawdown, meaning that our method can effectively manage the investment risk. In most years (62.5%), our portfolio makes a profit. It is only slightly worse than PXI in terms of 75.0% of profitable years. This shows the stability of our strategy.

Conclusions

We propose a deep learned-based investment strategy, which includes: (1) novel stock representation learning by deep CAE encoding of candlestick charts, (2) diversification through modularity optimisation based clustering and (3) portfolio construction by selecting the best Sharpe ratio stock in each cluster. Experimental results show: (a) our learned stock feature captures semantic information and (b) our portfolio outperforms the FTSE 100 index and many well-known funds in terms of total return, while providing better stability and lower risk than alternatives.

References

- [Amodei et al. 2016] Amodei, D.; Ananthanarayanan, S.; Anubhai, R.; et al. 2016. Deep speech 2: End-to-end speech recognition in english and mandarin. In *ICML*, 173–182.
- [Balsara and Zheng 2006] Balsara, N., and Zheng, L. 2006. Profiting from past winners and losers. *Journal of Asset Management*.
- [Barberis, Shleifer, and Vishny 1998] Barberis, N.; Shleifer, A.; and Vishny, R. 1998. A model of investor sentiment. *Journal of financial economics*.
- [Blondel et al. 2008] Blondel, V. D.; Guillaume, J.-L.; Lambiotte, R.; and Lefebvre, E. 2008. Fast unfolding of communities in large networks. *Journal of statistical mechanics: theory and experiment*.
- [Fama 1998] Fama, E. F. 1998. Market efficiency, long-term returns, and behavioral finance. *Journal of financial economics*.
- [Hakansson and Ziemba 1995] Hakansson, N. H., and Ziemba, W. T. 1995. Capital growth theory. *Handbooks in Operations Research and Management Science* 9:65–86.
- [He et al. 2016] He, K.; Zhang, X.; Ren, S.; and Sun, J. 2016. Deep residual learning for image recognition. In *CVPR*.
- [Huang et al. 2016] Huang, G.; Liu, Z.; Weinberger, K. Q.; and van der Maaten, L. 2016. Densely connected convolutional networks. *arXiv preprint arXiv:1608.06993*.
- [Jegadeesh and Titman 1993] Jegadeesh, N., and Titman, S. 1993. Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of finance* 48(1):65–91.
- [Keating and Shadwick 2002] Keating, C., and Shadwick, W. F. 2002. A universal performance measure. *Journal of performance measurement*.
- [Kelly 1956] Kelly, J. L. 1956. A new interpretation of information rate. *Bell Labs Technical Journal*.
- [Kingma and Welling 2013] Kingma, D. P., and Welling, M. 2013. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*.
- [Krizhevsky, Sutskever, and Hinton 2012] Krizhevsky, A.; Sutskever, I.; and Hinton, G. E. 2012. Imagenet classification with deep convolutional neural networks. In *NIPS*, 1097–1105.
- [LeCun et al. 1998] LeCun, Y.; Bottou, L.; Bengio, Y.; and Haffner, P. 1998. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*.
- [Lee et al. 2009a] Lee, H.; Grosse, R.; Ranganath, R.; and Ng, A. Y. 2009a. Convolutional deep belief networks for scalable unsupervised learning of hierarchical representations. In *ICML*. ACM.
- [Lee et al. 2009b] Lee, H.; Pham, P.; Largman, Y.; and Ng, A. Y. 2009b. Unsupervised feature learning for audio classification using convolutional deep belief networks. In *NIPS*, 1096–1104.
- [Lowe 1999] Lowe, D. G. 1999. Object recognition from local scale-invariant features. In *ICCV*. IEEE.
- [Maaten and Hinton 2008] Maaten, L. v. d., and Hinton, G. 2008. Visualizing data using t-sne. *Journal of Machine Learning Research* 2579–2605.
- [Markowitz and Todd 2000] Markowitz, H. M., and Todd, G. P. 2000. *Mean-variance analysis in portfolio choice and capital markets*, volume 66. John Wiley & Sons.
- [Markowitz 1952] Markowitz, H. 1952. Portfolio selection. *The journal of finance*.
- [Markowitz 1959] Markowitz, H. 1959. Portfolio selection: Efficient diversification of investments. cowles foundation monograph no. 16.
- [Masci et al. 2011] Masci, J.; Meier, U.; Cireşan, D.; and Schmidhuber, J. 2011. Stacked convolutional auto-encoders for hierarchical feature extraction. *Artificial Neural Networks and Machine Learning–ICANN 2011* 52–59.
- [Nanda, Mahanty, and Tiwari 2010] Nanda, S.; Mahanty, B.; and Tiwari, M. 2010. Clustering indian stock market data for portfolio management. *Expert Systems with Applications*.
- [Newman 2006] Newman, M. E. 2006. Modularity and community structure in networks. *Proceedings of the national academy of sciences*.
- [Onnela et al. 2003] Onnela, J.-P.; Chakraborti, A.; Kaski, K.; Kertesz, J.; and Kanto, A. 2003. Dynamics of market correlations: Taxonomy and portfolio analysis. *Physical Review E*.
- [Perronnin, Sánchez, and Mensink 2010] Perronnin, F.; Sánchez, J.; and Mensink, T. 2010. Improving the fisher kernel for large-scale image classification. *ECCV 2010* 143–156.
- [Povey et al. 2011] Povey, D.; Ghoshal, A.; Boulianne, G.; Burget, L.; Glembek, O.; Goel, N.; Hannemann, M.; Motlicek, P.; Qian, Y.; Schwarz, P.; et al. 2011. The kaldi speech recognition toolkit. In *IEEE 2011 workshop on automatic speech recognition and understanding*, number EPFL-CONF-192584. IEEE Signal Processing Society.
- [Rabiner and Juang 1993] Rabiner, L. R., and Juang, B.-H. 1993. Fundamentals of speech recognition.
- [Sharpe 1994] Sharpe, W. F. 1994. The sharpe ratio. *The journal of portfolio management*.
- [Silver et al. 2016] Silver, D.; Huang, A.; Maddison, C. J.; Guez, A.; Sifre, L.; Van Den Driessche, G.; Schrittwieser, J.; Antonoglou, I.; Panneershelvam, V.; Lanctot, M.; et al. 2016. Mastering the game of go with deep neural networks and tree search. *Nature*.
- [Simonyan and Zisserman 2014] Simonyan, K., and Zisserman, A. 2014. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*.
- [Szegedy et al. 2015] Szegedy, C.; Liu, W.; Jia, Y.; Sermanet, P.; Reed, S.; Anguelov, D.; Erhan, D.; Vanhoucke, V.; and Rabinovich, A. 2015. Going deeper with convolutions. In *CVPR*.
- [Tola et al. 2008] Tola, V.; Lillo, F.; Gallegati, M.; and Mantegna, R. N. 2008. Cluster analysis for portfolio optimization. *Journal of Economic Dynamics and Control*.

[Vincent et al. 2008] Vincent, P.; Larochelle, H.; Bengio, Y.; and Manzagol, P.-A. 2008. Extracting and composing robust features with denoising autoencoders. In *ICML*. ACM.